

SuperScalar

Lokwon Kim


Extracting Yet *More* Performance

- Two options:
 - Increase the depth of the pipeline to increase the clock rate – **superpipelining** (more details to come)
 - Fetch (and execute) more than one instructions at one time (expand every pipeline stage to accommodate multiple instructions) – **superscalar** (multiple-issue)

 - **Launching multiple instructions per stage allows the instruction execution rate, CPI, to be less than 1**
-

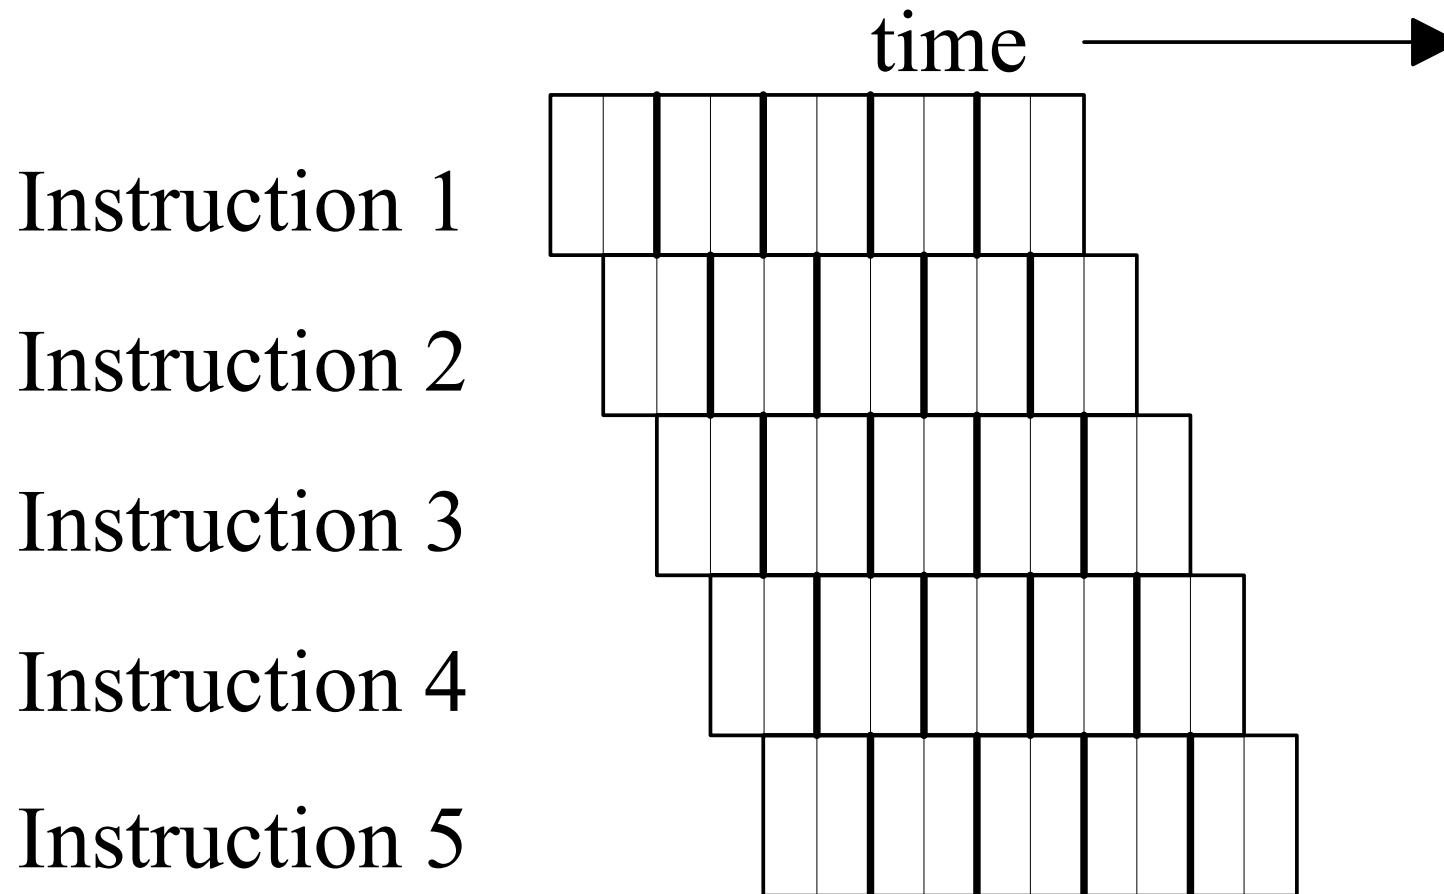
Instruction Pipelining

- ▣ **pipelining** – goal was to complete one instruction per clock cycle




Superpipelining

- **superpipelining** - Increase the depth of the pipeline to increase the clock rate



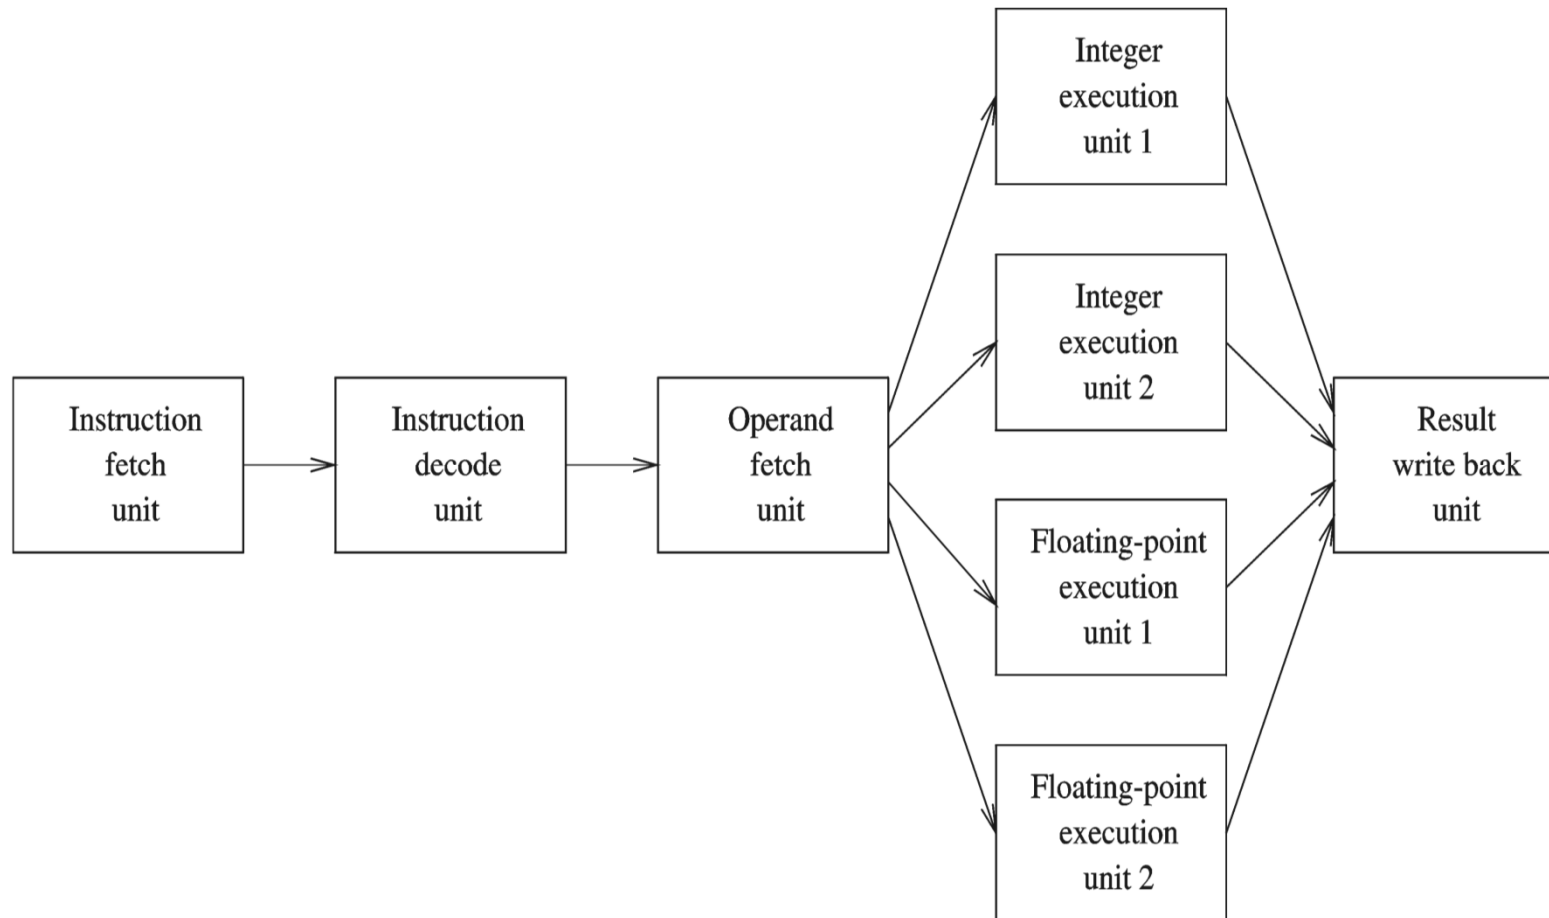
Superscalar – multiple-issue

- Fetch (and execute) more than one instructions at one time
(expand every pipeline stage to accommodate multiple instructions)



Superscalar – multiple-issue

- ❑ Fetch (and execute) more than one instructions at one time (expand every pipeline stage to accommodate multiple instructions)



Superpipelined Processors

- Increase the depth of the pipeline leading to shorter clock cycles (and more instructions “in flight” at one time)
 - Higher degree of superpipelining, (1) more forwarding/hazard hardware needed, (2) even more inevitable pipeline stalls (e.g. pipelined addition with data dependency), (3) gain of superpipelining is not linear (uneven balance of divided delay by superpipelining)

Superpipelined vs Superscalar(SS)

- **Superpipelined processors have longer instruction latency than the SS processors which can degrade performance in the presence of true dependencies**
 - **Superscalar processors are more susceptible to resource conflicts – but we can fix this with extra hardware !**
-

Multiple-Issue Processor Styles

- Static multiple-issue processors (aka **VLIW – very long instruction word architecture**)
 - Decisions on which instructions to execute simultaneously are being made statically (at compile time by the compiler)
 - E.g., Intel Itanium and Itanium 2 for the IA-64 ISA – EPIC (Explicit Parallel Instruction Computer)

 - **Dynamic multiple-issue processors** (aka **superscalar**)
 - Decisions on which instructions to execute simultaneously are being made dynamically (at run time by the hardware)
 - E.g., IBM Power 2, Pentium 4, MIPS R10K, HP PA 8500
-

Multiple-Issue Datapath (or SS) Responsibilities

- Must handle, with a combination of hardware and software fixes, the fundamental limitations of
 - Data dependencies – aka data hazards
 - Even more severe in SuperScalar (**Why?**).
 - Procedural dependencies – aka control hazards
 - Even more severe in SuperScalar (**Why?**).
 - Use dynamic branch prediction
 - Resource conflicts – aka structural hazards
 - Even more severe in SuperScalar (**Why?**).
 - Resource conflicts can be eliminated by duplicating the resource or by pipelining the resource
-

Out of Order (OoO) Execution

- **Instruction-issue** – initiate execution
 - **Instruction lookahead capability** – fetch, decode and issue instructions beyond the current instruction
- **Instruction-completion** – complete execution
 - **Processor lookahead capability** – complete issued instructions beyond the current instruction
- **Instruction-commit** – write back results to the RegFile or D\$ (i.e., change the machine state)

In-order issue with in-order completion

In-order issue with out-of-order completion

Out-of-order issue with out-of-order completion

Out-of-order issue with out-of-order completion and in-order commit

In-Order Issue with In-Order Completion

- **Simplest policy is to issue instructions in exact program order and to complete them in the same order they were fetched (i.e., in program order)**

- **Example:**

- Assume a pipelined processor (**3 stage pipeline**) that can fetch and decode **two** instructions per cycle, that has **three** functional units (a single cycle adder, a single cycle shifter, and a two cycle multiplier), and that can complete (and write back) **two** results per cycle
- And an instruction sequence with the following characteristics

I1 - needs two execute cycles (a multiply)

I2

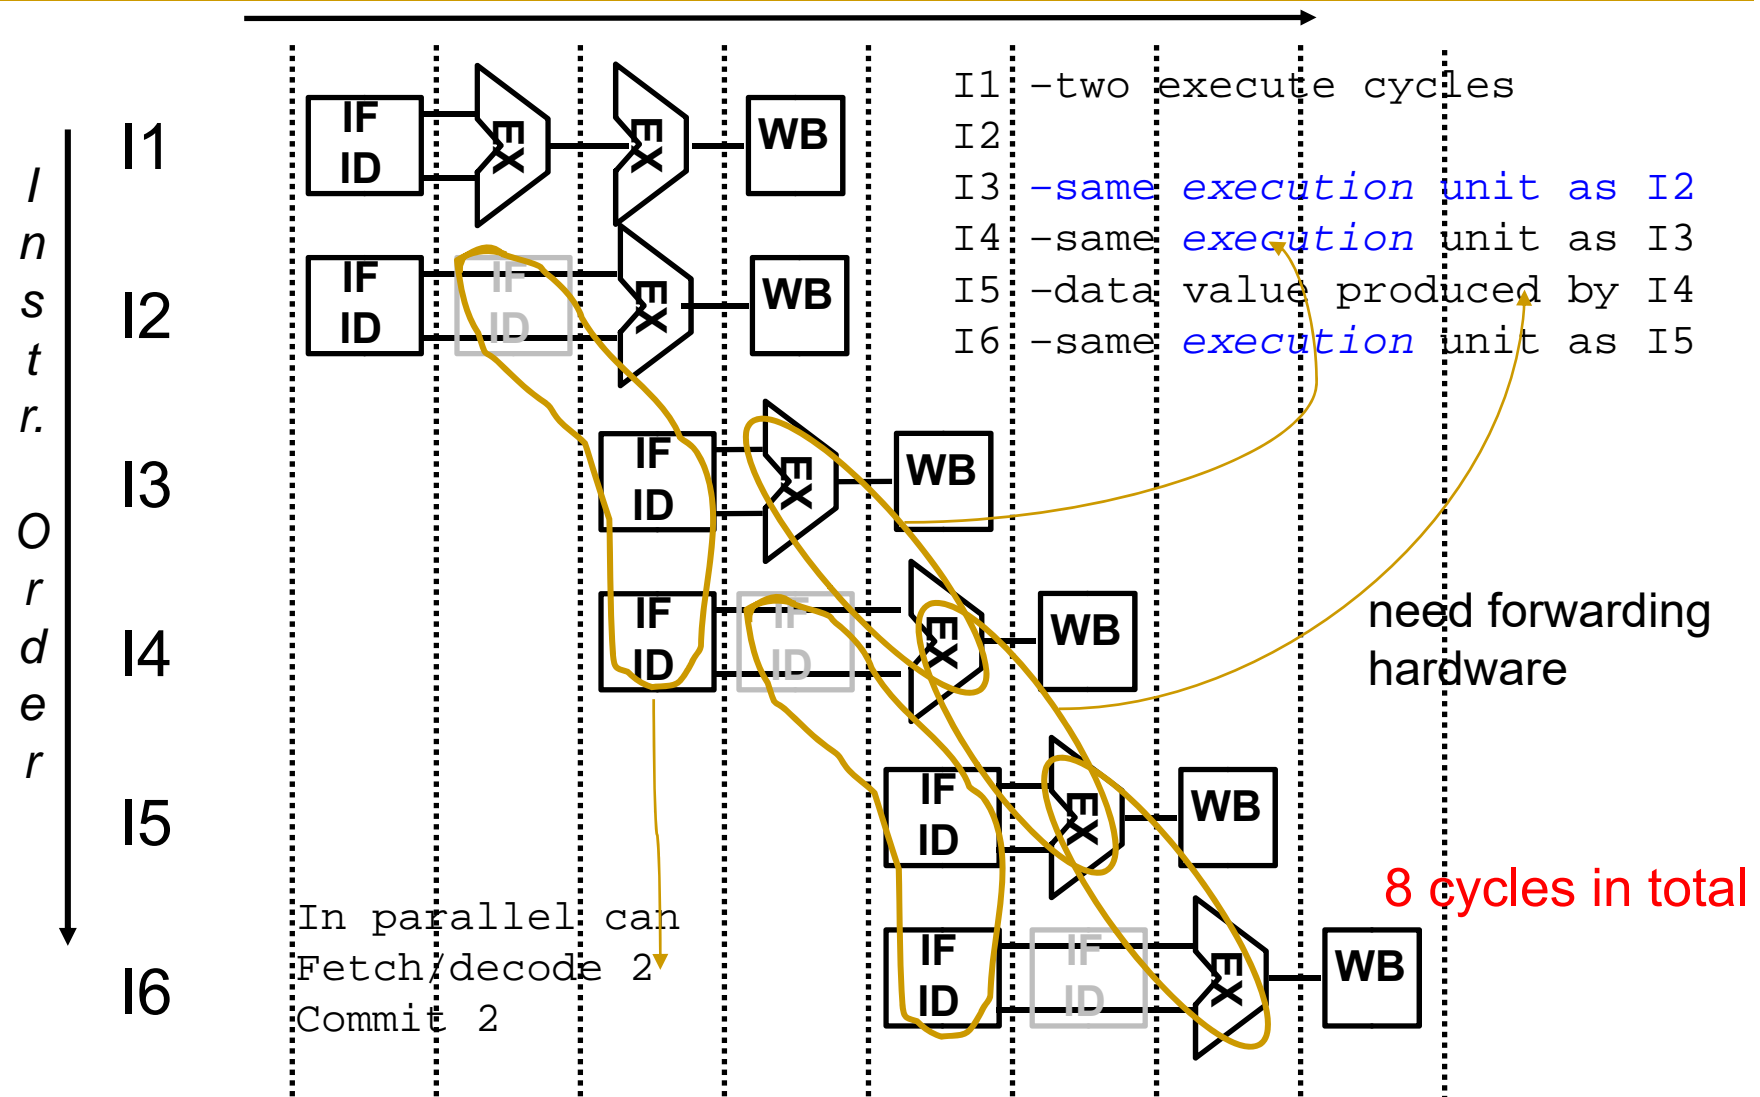
I3 - needs the same execution unit as I2

I4 - needs the same execution unit as I3

I5 - needs data value produced by I4

I6 - needs the same execution unit as I5

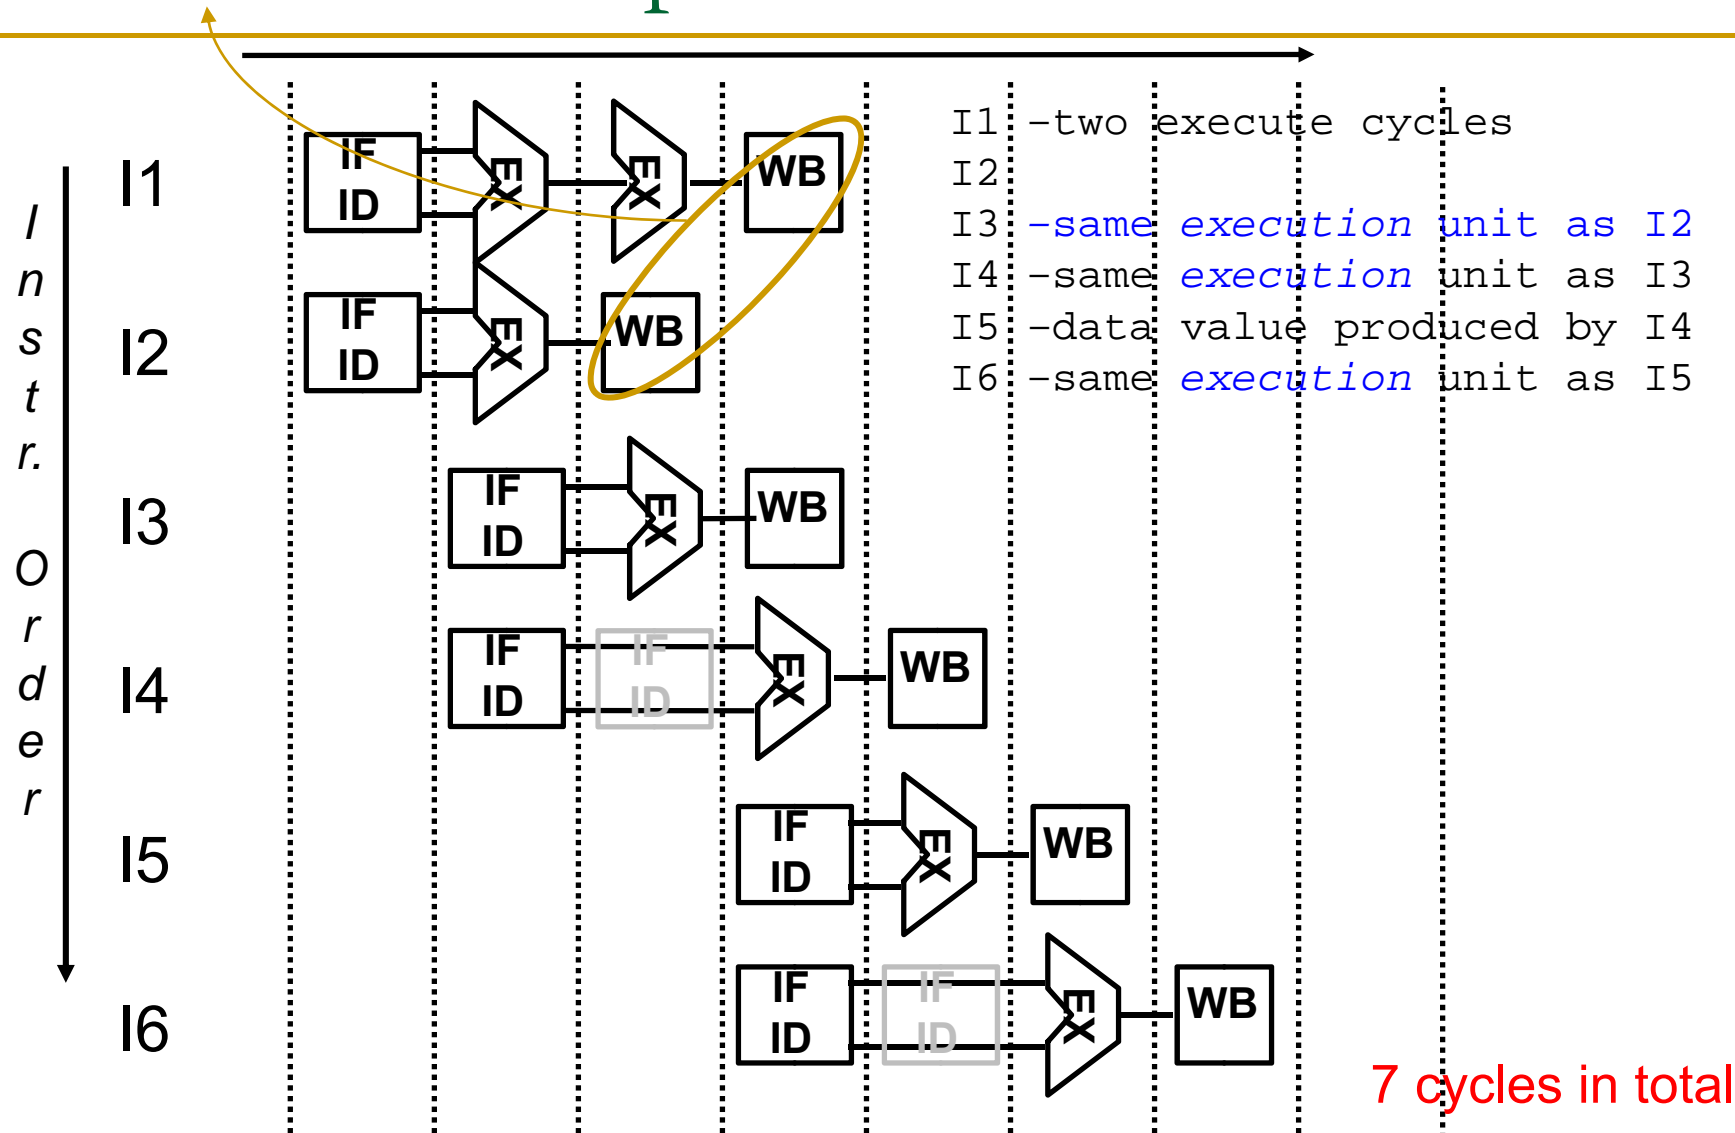
In-Order Issue, In-Order Completion Example



In-Order Issue with Out-of-Order Completion

- **With out-of-order completion, a later instruction may complete before a previous instruction**
 - Out-of-order completion is used in single-issue pipelined processors to improve the performance of long-latency operations such as divide
 - When using out-of-order completion, instruction is **stalled** when there is a resource conflict (e.g., for a functional unit) or when the instructions ready to issue need a result that has not yet been computed
-

IOI-OOC Example



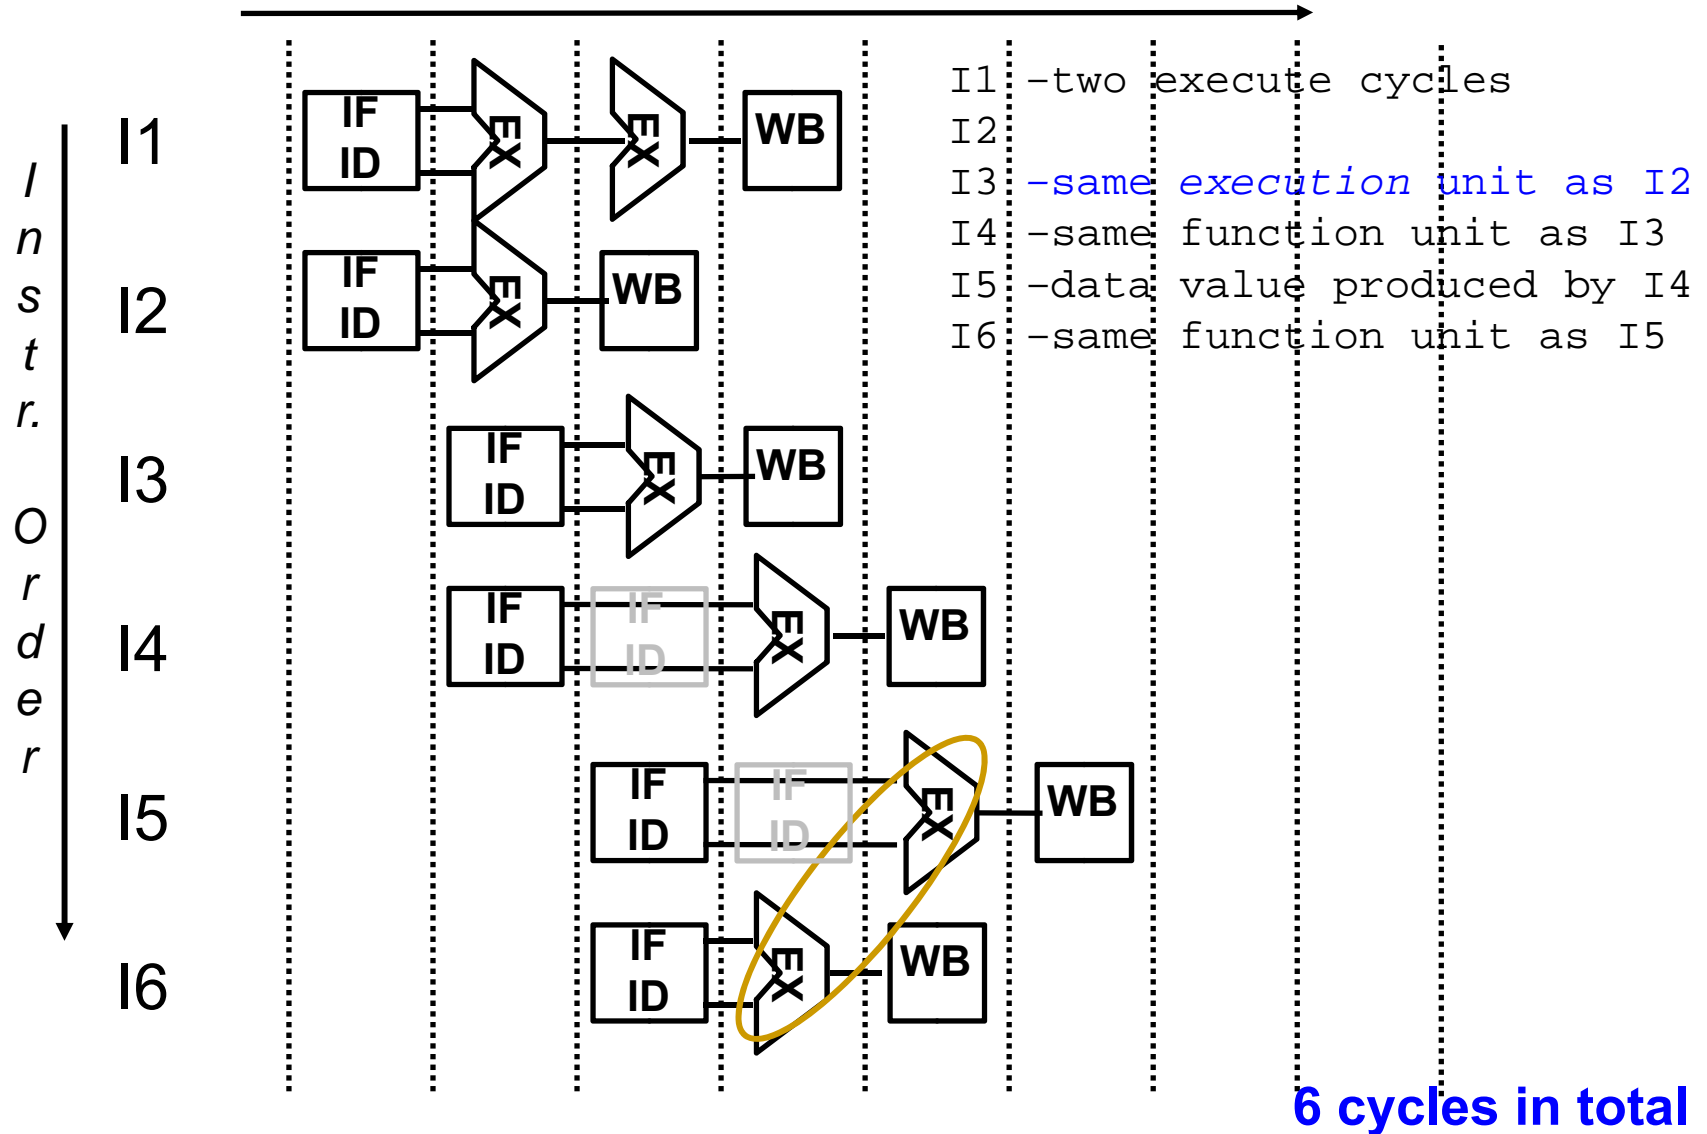
Handling Output Dependencies

- There is one more situation that stalls instruction issuing with IOI-OOC, assume
 - I1 – writes to R3
 - I2 – writes to R3
 - I5 – reads R3
 - If the I1 write occurs after the I2 write, then I5 reads an incorrect value for R3
 - I2 has an output dependency on I1 – write before write
 - The issuing of I2 would have to be stalled if its result might later be overwritten by an previous instruction (i.e., I1) that takes longer to complete – the stall happens before instruction issue
 - While IOI-OOC yields higher performance, it requires more dependency checking hardware
 - Dependency checking needed to resolve both read before write and write before write
-

Out-of-Order Issue with Out-of-Order Completion

- With in-order issue the processor stops decoding instructions whenever a decoded instruction has a resource conflict or a data dependency on an issued, but uncompleted instruction
 - The processor is not able to *look beyond* the conflicted instruction even though more downstream instructions might have no conflicts and thus be issueable
 - Fetch and decode instructions *beyond* the conflicted one, store them in an instruction buffer (as long as there's room), and flag those instructions in the buffer that don't have resource conflicts or data dependencies
 - Flagged instructions are then issued from the buffer without regard to their program order
 - **Issue means inserting instructions into Execution Unit.**
-

OOI-OOC Example



Antidependencies

- With OOI *also* have to deal with data antidependencies
 - when a later instruction (that completes earlier) produces a data value that destroys a data value used as a source in an earlier instruction (that issues later)

$\textcircled{R3} := R3 * R5$
 $R4 := R3 + 1$
 $\textcircled{R3} := R5 + 1$

True data dependency

Output dependency

Antidependency

- The constraint is similar to that of true data dependencies, except *reversed*
 - Instead of the later instruction using a value (not yet) produced by an earlier instruction (**read before write**), the later instruction produces a value that destroys a value that the earlier instruction (has not yet) used (**write before read**)
-

Dependencies Review

- Each of the three data dependencies
 - True data dependencies (read before write)
 - Antidependencies (write before read)
 - Output dependencies (write before write)

} storage conflicts

manifests itself through the use of registers (or other storage locations)
 - True dependencies represent the flow of data and information through a program
 - **Anti- and output dependencies arise because the limited number of registers mean that programmers reuse registers for different computations**
 - **When instructions are issued out-of-order, the correspondence between registers and values breaks down and the values *conflict* for registers**
-

Storage Conflicts and Register Renaming

- Storage conflicts can be reduced (or eliminated) by increasing or duplicating the troublesome resource
 - Provide additional registers that are used to reestablish the correspondence between registers and values
 - Allocated dynamically by the hardware in SS processors
- **Register renaming** – the processor renames the original register identifier in the instruction to a new register (one not in the visible register set)

$\textcircled{R3} := R3 * R5$		$R3b := R3 * R5$
$R4 := R3 + 1$	\Rightarrow	$R4 := R3b + 1$
$\textcircled{R3} := R5 + 1$		$R3c := R5 + 1$

Remove anti and output dependencies!

- The hardware that does renaming assigns a “replacement” register from a pool of free registers and releases it back to the pool when its value is superseded and there are no outstanding references to it
-

SuperScalar Summary

- Must handle, with a combination of hardware and software fixes, the fundamental limitations of
 - Data dependencies – aka data hazards
 - Using OOI-OOC makes even more severe in SuperScalar (**Why?**).
 - Procedural dependencies – aka control hazards
 - Even more severe in SuperScalar (**Why?**).
 - Use dynamic branch prediction
 - Resource conflicts – aka structural hazards
 - Even more severe in SuperScalar (**Why?**).
 - Resource conflicts can be eliminated by duplicating the resource or by pipelining the resource
-